

C++ Software Engineering Portfolio

Mahdi Tavakol — Postdoctoral Researcher, University of Oxford

Selected C++17/20 projects demonstrating modern design, numerical algorithms, and parallel computing on Linux.

Focus Areas:

- ❑ RAII, Eigen, and exception-safe design
- ❑ Object Oriented code design
- ❑ OpenMP and MPI parallelism
- ❑ CMake and Git
- ❑ High-performance numerical algorithms and simulation frameworks



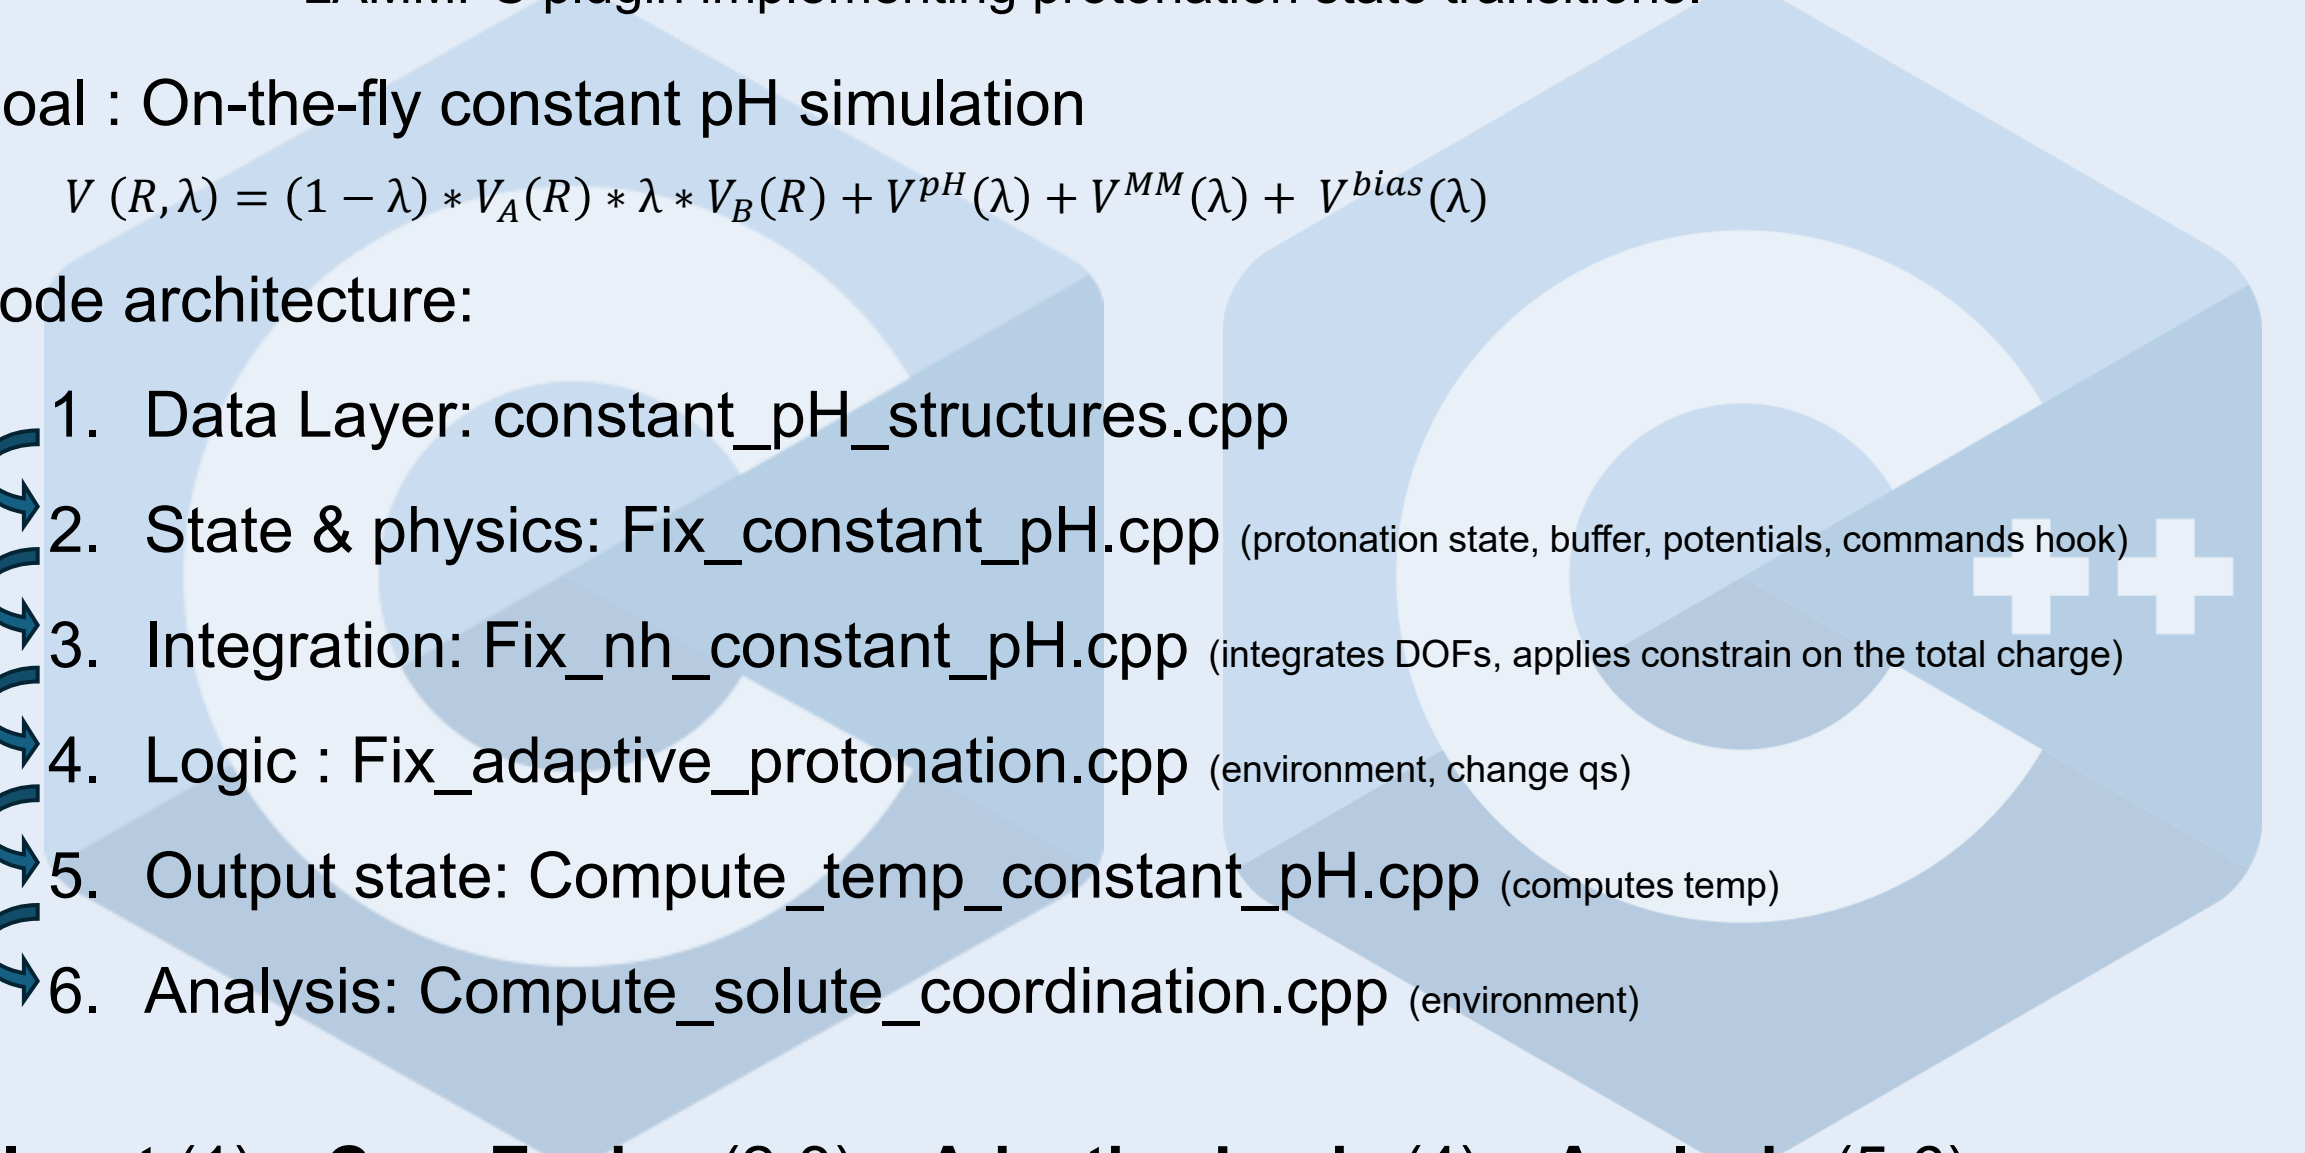
Project #1: Constant – pH simulation

LAMMPS plugin implementing protonation state transitions.

Goal : On-the-fly constant pH simulation

$$V(R, \lambda) = (1 - \lambda) * V_A(R) * \lambda * V_B(R) + V^{pH}(\lambda) + V^{MM}(\lambda) + V^{bias}(\lambda)$$

Code architecture:

- 
1. Data Layer: `constant_pH_structures.cpp`
 2. State & physics: `Fix_constant_pH.cpp` (protonation state, buffer, potentials, commands hook)
 3. Integration: `Fix_nh_constant_pH.cpp` (integrates DOFs, applies constrain on the total charge)
 4. Logic : `Fix_adaptive_protonation.cpp` (environment, change qs)
 5. Output state: `Compute_temp_constant_pH.cpp` (computes temp)
 6. Analysis: `Compute_solute_coordination.cpp` (environment)

Input (1) – Core Engine (2-3) – Adaptive Logic (4) – Analysis (5-6)

Project #1: Constant – pH simulation

constant_pH_structures.cpp

RAII + smart_ptr

```
constant_pH_structures(LAMMPS *lmp, const string &fileName1, const string &fileName2;

~constant_pH_structures();
constant_pH_structures(const constant_pH_structures &rhs) = delete;
constant_pH_structures &operator=(const constant_pH_structures &rhs) = delete;
constant_pH_structures(constant_pH_structures &&rhs) = delete;
constant_pH_structures &operator=(constant_pH_structures &&rhs) = delete;
```

```
void read_pH_structure_files();
```

```
private:
ifstream pHStructureFile1, pHStructureFile2;
```

```
/* These variables will be accessed by fix_adaptive_protonation and fix_constant_pH*/
```

```
std::unique_ptr<int[]> typePerProtMol;
std::unique_ptr<int[]> protonable;
int pmnstructures1, pmnstructures2;
int pHnTypes1, pHnTypes2;
```

```
if (flags & ADAPTIVE) {
```

```
molids = std::make_unique<int[]>(n_lambdas);
// get protonable molids and the molids that are compatible with the unique_ptr
fix_adaptive_protonation->get_protonable_molids(molids.get());
}
```

fix_constant_pH.cpp: keeping lambdas from the prev step

```
int to = 0;
if (molids_prev) {
for (int i = 0; i < n_lambdas; i++) {
auto iter = std::find(molids_prev.get(), molids_prev.get()+n_lambdas_prev, molids[i]);
if (iter != molids_prev.get()+n_lambdas_prev) {
int from = std::distance(molids_prev.get(), iter);
for (int j = 0; j < 3; j++) {
lambdas[to][j] = lambdas_prev[from][j];
v_lambdas[to][j] = v_lambdas_prev[from][j];
a_lambdas[to][j] = a_lambdas_prev[from][j];
m_lambdas[to][j] = m_lambdas_prev[from][j];
}
H_lambdas[to] = H_lambdas_prev[from];
to++;
}
}
```

fix_constant_pH:
Getter and Setter Funcs

```
// Functions for accessing or resetting the lambda dynamics parameters
void return_n_params(int &n_params) const;
void return_params(double **const _x_lambdas, double **const _v_lambdas,
double **const _a_lambdas, double **const _m_lambdas) const;
void reset_params(double **const _x_lambdas, double **const _v_lambdas, double **const _a_lambdas,
double **const _m_lambdas, const int mod = 1);
void reset_params(double **const _x_lambdas, double **const _v_lambdas, double **const _a_lambdas,
double **const _m_lambdas);

reset_params(_x_lambdas, _v_lambdas, _a_lambdas, _m_lambdas, 1);

void return_H_lambdas(double *H_lambdas) const;
void return_T_lambda(double &T_lambda, int component = 2);

// Functions to return the buffer parameters
void return_buff_params(double &x_lambda_buff, double &v_lambda_buff, double &a_lambda_buff,
double &m_lambda_buff, int &N_buff) const;
void reset_buff_params(const double _x_lambda_buff, const double _v_lambda_buff,
const double _a_lambda_buff, const double _m_lambda_buff, const int mod = 1);
void reset_buff_params(const double _x_lambda_buff, const double _v_lambda_buff,
const double _a_lambda_buff, const double _m_lambda_buff);

reset_buff_params(_x_lambda_buff, _v_lambda_buff, _a_lambda_buff, _m_lambda_buff, 1);

// Function to set the charges based on the lambdas and lambda_buff values
void reset_qs();

void calculate_Hs();
void check_num_OWs_HWs();
void read_pH_structure_files();
void read_commands_file();
void restore_epsilon();
```

fix_constant_pH.cpp:
MPI_comm

```
std::unique_ptr<double[]> Us;
std::unique_ptr<double[]> dUs;

// parameter for shifting the minima of the potential near lambda = 0 and lambda = 1
double mu;

// Random number seed for the initialization of initial v_lambdas
double random_number_seed;

// Buffer potential parameters
double a_buff, b_buff, s_buff, m_buff, w_buff, r_buff, d_buff, k_buff, h_buff;

// Lambda arrays
double **lambdas, **v_lambdas, **a_lambdas, **m_lambdas, *H_lambdas;
double T_lambdas[3];
```

```
void FixConstantPH::grow_arrays(int nmax_new)
{
double *new_vector_atom = new double[nmax_new];
int keep = std::min(nmax, nmax_new);
if (keep > 0) std::copy(vector_atom, vector_atom+keep, new_vector_atom);
if (keep < nmax_new) std::fill(new_vector_atom+keep, new_vector_atom+nmax_new, 0);
delete [] vector_atom;
vector_atom = new_vector_atom;
nmax = nmax_new;
}
```

```
if (lambda_thermostat_type == LAMBDA_ANDERSEN && comm->me == 0) { ... }
else if (lambda_thermostat_type == LAMBDA_BUSSI && comm->me == 0) { ... }
else if (lambda_thermostat_type == LAMBDA_NOSEHOOVER && comm->me == 0) { ... }
```

```
if (n_lambdas == 0)
return;
```

fix_nh_constant_pH.cpp:
deterministic

```
// v_lambdas[0] is the location of the contiguous memory allocated
// for the double ** v_lambdas
```

```
MPI_Bcast(v_lambdas[0], n_lambdas*3, MPI_DOUBLE, 0, world);
```

```
if (lambda_integration_flags & BUFFER)
MPI_Bcast(&v_lambda_buff, 1, MPI_DOUBLE, 0, world);
```

```
for (int i = 0; i < n_lambdas; i++) {
v_cm += v_lambdas[i][0]*mols_charge_change;
}
```

Modern C++ 17 – RAII – MPI – Deterministic Integration – Modular Design

Project #2: C++17 Plugin Architecture for Adaptive Biasing (LAMMPS Extension)

Adaptive bias potential framework for free-energy surface exploration.

Goal : Apply adaptive biasing forces to escape metastable states

Code architecture:

1. Parent abstract class: `compute_diff_atom.cpp` (pure abstract class)
 - a) `compute_entropy_atom.h` (user defined CVs)
 - b) `compute_cparam_atom.h` (user defined CVs)
 - c) `compute_enthalpy_atom.h` (user defined CVs)
 - d) `compute_q6_smooth_atom.h` (complex math, many ghost communications, on-the-fly computation due to large memory demand)
2. State & physics: `Fix_metaAR.h` (Integration logic, file I/O, history management, bias update scheme)
3. Integration: `fix_nvt.h` (integrates DOFs, applies constraint on the total charge)

Input (1) – Core Engine (2-3) – Adaptive Logic (2) – Analysis (2)

Project #2: C++17 Plugin Architecture for Adaptive Biasing (LAMMPS Extension)

constant_diff_atom.cpp

Inheritance + modern C++ 17

```
class ComputeDiffAtom : public Compute {  
public:  
    ComputeDiffAtom(class LAMMPS *, int, char **);  
    ~ComputeDiffAtom() override;  
    virtual void init() override;  
    void init_list(int, class NeighList *) override;  
    double compute_scalar() override;  
    void compute_peratom() override;
```

```
static constexpr int val_col = 0;  
static constexpr int diff_x_col = 1;  
static constexpr int diff_y_col = 2;  
static constexpr int diff_z_col = 3;
```

compute_q6_smooth_atom.cpp

Input parsing

```
int iarg = 5;  
while (true) {  
    if (strcmp(arg[iarg], "no_diff") == 0) {  
        if (comm->me == 0) error->warning(FLError, "Turning off the diffs, this con  
        iarg++;  
    } else if (strcmp(arg[iarg], "phi") == 0) {  
        mode = (mode & ~N_MODE) | PHI_MODE;  
        iarg++;  
    } else if (strcmp(arg[iarg], "S0_off") == 0) {  
        switch_flag &= ~S0_SW;  
        iarg++;  
    } else if (strcmp(arg[iarg], "S1_off") == 0) {  
        switch_flag &= ~S1_SW;  
        iarg++;  
    } else if (strcmp(arg[iarg], "S2_off") == 0) {  
        switch_flag &= ~S2_SW;  
        iarg++;  
    } else if (strcmp(arg[iarg], "S3_off") == 0) {  
        switch_flag &= ~S3_SW;  
        iarg++;  
    } else if (strcmp(arg[iarg], "S0") == 0) {  
        if (iarg + 1 >= narg) error->all(FLError, "Missing parameters after S0");  
        threshold = utils::numeric(FLError, arg[iarg + 1], false, 1mp);  
        iarg++;
```

fix_metaAR.cpp: testing

```
for (int j = 0; j < binsCV2; j++) {  
    double Bt = 0.0;  
    bias[i][j] = 0.0;  
    for (int k = 0; k < n - 1; k++) {  
        double A = 0.5 *  
            ((locCV1[i] - CV1[k]) * (locCV1[i] - CV1[k]) / (sigmaCV1 * sigmaCV1) +  
             (locCV2[j] - CV2[k]) * (locCV2[j] - CV2[k]) / (sigmaCV2 * sigmaCV2));  
        Bt += 0.592 * omega * exp(-biasHistory[k] / (boltz * deltaT)) * exp(-A);  
    }  
    bias[i][j] = Bt;
```

```
class FixMetaARTest : public LAMMPTTest {  
protected:  
    void SetUp() override {  
        testbinary = "FixMetaARTest";  
        LAMMPTTest::SetUp();  
    }
```

compute_q6_smooth_atom:

$$N[i] = s2 \left(\sum_j s1(b_{ij}) * s0(r_{ij}) \right)$$

$$\Phi[i] = \sum \sum s3(r_{ij}) * N[i] * N[j]$$

$$bij = \sum_{m=-6}^{m=6} q6[i][m] * q6[j][m]$$

$$\Phi' = N[i] * \left(\sum \frac{dNi}{drj} \right) + \frac{dNi}{dri} * \left(\sum Nj \right)$$

1. Central atom q6m, diffq6m and coeffs
2. Fill ghost atoms q6m, diffq6m, s2, ds2 and coeffs
3. $dcij/dri = qj*dqi/dri(A) + qi*dqj/dri(B)$ (We do not have the B)
 1. Fill A in for the atom i and fill B for the atom j in h[j]
 2. Add from the h calculated by the ghost atoms
4. Add the diff related to s2 and s0 (distance)
5. We have Ni and diffNi/diffri
6. Send Ni to ghost atoms
7. Calculate the dNj/dri on-the-fly from qj and coeffs

RAII-safe memory allocation (LAMMPS Memory API)- Cxxpr – Virtual overrides
– C++ smart pointers – On-the-fly calculations to avoid heap exhaustion

Project #3: Free energy Part A: C++ Thermodynamic Integration module

Custom C++ LAMMPS compute for thermodynamic-integration (TI) energy derivatives.

Code architecture:

1. Parent abstract class: compute.h (pure abstract class)
compute_thermo_integ.h

```
template <int direction>
void ComputeThermoInteg::backup_restore_qfev()
{
    int i;

    int nall = atom->nlocal + atom->nghost;
    int natom = atom->nlocal;
    if (force->newton || force->kspace->tip4pflag)

template <int direction>
void ComputeThermoInteg::forward_reverse_copy(d
{
    if (direction == 1) a = b;
    if (direction == -1) b = a;
}
```

templates

```
std::map<std::string, std::string> pair_params;

pair_params["lj/cut/soft/omp"] = "lambda";
pair_params["lj/cut/coul/cut/soft/gpu"] = "lambda";
pair_params["lj/cut/coul/cut/soft/omp"] = "lambda";
pair_params["lj/cut/coul/long/soft"] = "lambda";
pair_params["lj/cut/coul/long/soft/gpu"] = "lambda";
pair_params["lj/cut/coul/long/soft/omp"] = "lambda";
pair_params["lj/cut/tip4p/long/soft"] = "lambda";
pair_params["lj/cut/tip4p/long/soft/omp"] = "lambda";
pair_params["lj/charmm/coul/long/soft"] = "lambda";
pair_params["lj/charmm/coul/long/soft/omp"] = "lambda";
pair_params["lj/class2/soft"] = "lambda";

std::maps
```

```
// ----- Fixture 1 -----
class PeVsUCTest : public LAMMPSSTest {
protected:
    void SetUp() override {
        testbinary = "PeVsUCTest";
        LAMMPSSTest::SetUp();

        const double RefB_All = get_equal(lmp, "RefB_All");
        const double RefB_HAp = get_equal(lmp, "RefB_HAp");

        // ----- Compare -----
        constexpr double tol = 1e-3;
        EXPECT_NEAR(UA_All, RefA_All, tol);
        EXPECT_NEAR(UA_HAp, RefA_HAp, tol);
        EXPECT_NEAR(UB_All, RefB_All, tol);
        EXPECT_NEAR(UB_HAp, RefB_HAp, tol);
    }
};

testing
```

```
/* -----
ComputeThermoInteg::~ComputeThermoInteg()
{
    deallocate_storage();
    memory->destroy(epsilon_inits);
    memory->destroy(energy_peratom);

    void ComputeThermoInteg::deallocate_storage()
    {
        memory->destroy(vector_atom);
        memory->destroy(q_orig);
        memory->destroy(f_orig);

RAII implementation for
double**
(mdspan was not available yet!)
```

```
q_local += q[i]; MPI (Distributed parallelisation)
MPI_Allreduce(&q_local, &q_total, 1, MPI_DOUBLE, MPI_SUM, world);
```

Verified, reproducible, and parallel-safe
C++ compute module integrated with
LAMMPS core

Project #3: Free energy Part B: C++ and Python Pipeline for high throughput data

```
#pragma omp parallel for
```

OpenMP

```
for (int i = 1; i <= numFolders * numSims; i++) {
```

```
string command = "JAX_ENABLE_X64=True ./BAR-v03.py " +  
to_string(numData) + " " + to_string(i);
```

python
from
C++

```
array<char, 128> buffer;
```

```
string resultString;
```

```
unique_ptr<FILE, decltype(&pclose)> pipe(popen(command.c_str(), "r"),  
pclose);
```

```
#pragma omp critical
```

Thread safety

```
std::cerr << "Failed to run command: " << command << "\n";  
continue;  
}
```

Making a C command RAI
compatible

```
template <typename... Args>
```

```
JarCalculator::JarCalculator(const string &logName_, const int &numFolders_,  
const int &numSims_, const int &numData_,  
Args &&...args)  
: logName(logName_), numFolders(numFolders_), numSims(numSims_),  
numData(numData_) {  
static_assert((std::is_convertible_v<Args, std::string> && ...),  
"All arguments must be convertible to std::string");  
vector<string> foldersTmp{std::forward<Args>(args)...};
```

Variadic Templates

```
double sq_sum =
```

STL algorithms

```
std::accumulate(gBarShuffleI, gBarShuffleI + numShuffles, 0.0,  
[gBar_meanI, etol](double acc, double val) {  
if (std::abs(val) >= etol) {  
double diff = val - gBar_meanI;  
return acc + diff * diff;  
} else  
return acc;  
});
```

```
std::random_device rd;
```

```
std::mt19937_64 g{rd()};
```

```
for (int i = 0; i < numShuffles; i++) {  
auto firstIndex = shuffles.begin() + i * numFolders * numSims;  
auto lastIndex = firstIndex + numFolders * numSims;
```

```
std::iota(firstIndex, lastIndex, 0);  
std::shuffle(firstIndex, lastIndex, g);
```

```
// I used resize on purpose so the initial values are set to zero!  
step.resize(numDirs * numData); // forward-reverse  
z1.resize(numSims * numData * numFolders *  
numDirs); // forward-reverse + two folders  
z2.resize(numSims * numData * numFolders *  
numDirs); // forward-reverse + two folders  
for (int i = 0; i < numFolders * numSims; i++) {  
int indx =  
i + numFolders * numSims * j; // (i + z * numSims, j + k * numData )  
datai = data[indx];
```

N-D vector + contiguous
memory

Parallel, memory-efficient, and language-integrated C++ pipeline for high-throughput free-energy computation

Project #4: Mandelbrot Parallel Renderer

A Modular C++ Framework for Parallel Simulation and Rendering

(1) Main.cpp

— (2) Runner/ (run_mandelbrot_timing, run_mandelbrot_animation)
 └─ Factory-like controller

— (3) Algorithm/
 └─ mandelbrot.h
 └─ mandelbrot_xmesh(_inner/outerloop).h (inheritance)
 └─ mandelbrot_xmesh(_inner/outerloop).h

— (4) Array/
 └─ array.h (Base interface)
 └─ array_c / array_cpp / array_modern / array_mdspan
 └─ array_allocator (RAII factory)

— (5) Numerical/
 └─ complex.h (templates)

— (6) Plotting/
 └─ Mandelbrot_plot_frames.py

Input (1) – Control Layer (2) – Core Engine (3) – Memory & Data Layer (4) – Numerical Library (5) – Post-processing (6)

Project #4: Mandelbrot Parallel Renderer

```
namespace Array_NS {  
    class array_allocator  
    {  
    public:  
        array_allocator(allocation_mode mode_, const allocation_major& major_,  
            const int& n_xs_, const int& n_ys_, const std::string& output_name_)  
            : mode{ mode_ }, major{ major_ },  
            n_xs{ n_xs_ }, n_ys{ n_ys_ },  
            output_name{ output_name_ }  
        {}  
  
        if (major == allocation_major::X_MAJOR)  
        {  
            switch (mode) {  
            case allocation_mode::C:  
                array_ptr = std::make_unique<array_c<allocation_major::X_MAJOR>>(n_xs, n_ys, output_name);  
                break;  
            case allocation_mode::CPP:  
                array_ptr = std::make_unique<array_cpp<allocation_major::X_MAJOR>>(n_xs, n_ys, output_name);  
                break;  
            case allocation_mode::MODERN:  
                array_ptr = std::make_unique<array_modern<allocation_major::X_MAJOR>>(n_xs, n_ys, output_name);  
                break;  
            }  
        }  
  
        double ara = 0.0;  
#pragma omp parallel reduction(+:ara) // default(none) shared(n_xs,x_min,x_max,y_min,y_max,y_si  
{  
    int thread_id = omp_get_thread_num();  
    if (thread_id == 0)  
        std::cout << "Using " << omp_get_num_threads() << " omp threads" << std::endl;  
    int first = first_ranges[thread_id];  
    int last = last_ranges[thread_id];  
  
    for (int j = 0; j < n_ys; j++)  
    {  
        for (int i = first; i < last; i++)  
        {  
            complex min(static_cast<double> (this->x_min), static_cast<double> (this->y_min)  
            double _i = static_cast<double> (i % this->n_xs);  
            double _j = static_cast<double> (j);  
            _i = static_cast<double> (( _i + this->n_xs / _scale) / _scale);  
            _j = static_cast<double> (( _j + this->n_ys / _scale) / _scale);  
  
            double& operator()(int x, int y) override {  
                if (bounds_check_flag)  
                    bounds_check(x, y); // Check bounds before accessing  
                if (alloc_major == allocation_major::X_MAJOR) {  
                    auto mdspn = std::mdspan(data.data(), n_xs, n_ys);  
                    return mdspn[x, y];  
                }  
            }  
        }  
    }  
}
```

namespace +
Factory pattern

Parallelism +
Threads safety

C++23 style

```
namespace Array_NS {  
    template<allocation_major alloc_major> <T> Provide sample template arguments for IntelliSense  
    class array_cpp : public array  
    {  
    public:  
        array_cpp(const int& _n_xs, const int& _n_ys, const std::string& _output) :  
            array{ _n_xs, _n_ys, _output }  
        {}  
  
        array_cpp(const array_cpp& _in) = delete;  
        array_cpp& operator=(const array_cpp& _in) = delete;  
        array_cpp(array_cpp& _in) = default;  
        array_cpp& operator=(array_cpp& _in) = default;  
        ~array_cpp() {  
            deallocate();  
        }  
  
        double& operator()(int x, int y) override {  
            if (bounds_check_flag)  
                bounds_check(x, y); // Check bounds before accessing  
            if (alloc_major == allocation_major::X_MAJOR)  
            {  
                double* temp = (double*)malloc(n_xs * n_ys * sizeof(double));  
                if (alloc_major == allocation_major::X_MAJOR) {  
                    data = (double**)malloc(n_ys * sizeof(double*));  
                    for (int i = 0; i < n_xs; i++)  
                        data[i] = &temp[i * n_ys];  
                }  
                else if (alloc_major == allocation_major::Y_MAJOR) {  
                    data = (double**)malloc(n_ys * sizeof(double*));  
                    for (int i = 0; i < n_ys; i++)  
                        data[i] = &temp[i * n_xs];  
                }  
            }  
            void deallocate() override {  
                if (data)  
                {  
                    free(data[0]);  
                    free(data);  
                }  
            }  
        }  
  
        int main([[maybe_unused]] int argc, [[maybe_unused]] char** argv)  
        {  
            using Mandelbrot_NS::bounds;  
            using Runner_NS::run_mandelbrot_timing;  
            try {  
                auto run_mandelbrot_ptr = std::make_unique<run_mandelbrot_timing>(bnds, x_size, y_size)  
                run_mandelbrot_ptr->run();  
            }  
            catch (std::bad_alloc& ae) {  
                std::cerr << "C++ allocation failed:" << ae.what() << std::endl;  
                return EXIT_FAILURE;  
            }  
            catch (std::exception& ae) {  
                std::cerr << "Exception: " << ae.what() << std::endl;  
            }  
        }  
    }  
}
```

Template +
Operator overloading

Contiguous memory
allocation for cache
locality

Exception Safety +
RAII Smart Pointers

Project #5: Neural Network:

A Modular C++ Framework for Training Feed-Forward Nets

(1) Input/Configuration

- ❖ main.cpp
- ❖ InputArgs.h --- Parsing CLI
- ❖ InputFile.h --- CSV reader
- ❖ InputFileMPI.h --- Parallel CSV reader

(2) Data Processing & Logging

- ❖ Scaler.h --- MinMaxScaler,...
- ❖ Logger.h --- log to cout, log file, multi-stream

(3) Core Engine

- ❖ NeuralNetwork.h --- Core training/Feed-forward
- ❖ NeuralNetworkMPI.h --- MPI accelerator
- ❖ NeuralNetworkOpenMP.h --- OpenMP accelerator

(4) Model Components

- ❖ Layer.h --- Layers
- ❖ Activation.h --- Relu/Sigmoid/Tanh
- ❖ Dropout.h --- Regularization

(5) Numerical Library

- ❖ Eigen

(6) Optimization

- ❖ Loss.h --- MSE, MAE, Huber
- ❖ Optimizers.h --- SGD, RMSProp, ...

(7) Builds and tests

- ❖ CMakeLists.txt --- Debug/release, MPI/OpenMP
- ❖ tests/ --- regression test with catch

Input (1) → Data (2) → Core/Model (3-4) → Math/Opt (5-6) → Builds/tests (7)

Project #5: Neural Network

InputArgs : Easy for mock in tests

MPI/OpenMPI accelerator through inheritance

RAII resource management + Eigen

Eigen

Transpose: Data are in columns and Eigen is col-major → Cache efficiency in parallel

Multilevel logging

Multistream logger

Summary and Discussions

- ❖ Designed and implemented multiple C++17/20 scientific frameworks using RAII, Eigen, OpenMP, MPI and CUDA.
- ❖ Focused on numerical stability, deterministic parallelism and cache-efficient computation.
- ❖ Modular, testable and cross-compiler code with unit testing.
- ❖ Applied design patterns (Factory, Interface, Strategy) for extensibility.